

AI Critique

AI was useful for expanding checklist coverage, identifying suspicious code paths, automating repetitive browser actions, and reconciling evidence across several documents. However, its first weakness was treating plausible source-level patterns as if they were established runtime facts. Several early defect descriptions assumed controls or screens behaved in a particular way before the running interface had been inspected. The final process therefore separated “candidate hypothesis” from Passed, Failed, and “verified defect,” and only the latter categories were supported by repeated execution.

The automation work exposed a second limitation: an AI-generated test can fail for reasons unrelated to the SUT. Initial runs stopped because of incorrect state setup, asynchronous browser-dialog handling, premature navigation assertions, and an imprecise revenue parser. Those failures could easily have been misreported as product defects. Resetting the database and rerunning the complete suite after each harness correction was essential. The AI also initially used Playwright Chromium while the requirement explicitly named Chrome; human review changed the final run to the installed Google Chrome channel.

Finally, AI could not establish Firefox compatibility, physical-device behavior, participant identity, consent, frustration, or SUS responses. Firefox launched but could not create a usable page in the available environment, so the honest outcome was Blocked, not a fabricated pass. The main lesson is that AI accelerates analysis and execution but does not replace evidentiary judgment. Reliable collaboration requires explicit evidence thresholds, deterministic setup, repeated runs, visual review, and transparent reporting of anything the available environment cannot prove.