

AI Critique

AI accelerated the external-JSON design, Playwright configuration, and first implementation of 51 cases for FR-06, FR-10, and FR-12, but its output still required technical and evidential review. The clearest generated-code error was the original FR-10 order-row locator: substring matching for `#2` could also select `#20`, creating a false failure. Human review replaced it with an exact table-cell match and repeated all three FR-10 browser runs; `FR10-02` then passed in the selected reports. AI also could not infer that Firefox failures were infrastructural merely from a failed run. Headless Firefox failed before `newPage()`, while headful smoke testing succeeded but the full matrix still produced 22 intermittent page-fixture timeouts. Treating those as product defects would have distorted the defect count, so they were separated from the 52 assertion failures. The final audit also exposed subtler gaps that execution alone did not catch: invalid FR-06 quantities are checked through native validity without proving that cart state remains unchanged after an Add attempt, and FR-12 cleanup requests are not followed by assertions that deletion succeeded. These omissions occurred because the model optimized for broad case coverage and runnable structure while lacking a complete state oracle, reliable browser-runtime context, and an explicit requirement to prove cleanup postconditions. The principle learned is that AI-generated tests are hypotheses, not evidence. Every locator, state transition, cleanup action, and failure classification needs human review against the requirement; every correction needs a rerun; and only attributable artifacts may support a pass, product defect, or environment-failure claim.

Word count: 246 words (paragraph only; counted by tokenising letters/numbers and internal apostrophes/hyphens).